




Data Management

Sahar Khanlari - 2107563

Marco Natale – 1929854

A.A. 2024/2025




Introduction

- SQL vs NoSQL (PostgreSQL vs MongoDB)
- Goals:
 - Interface Usability
 - Flexibility
 - Scalability
 - Performance Queries and Comparison

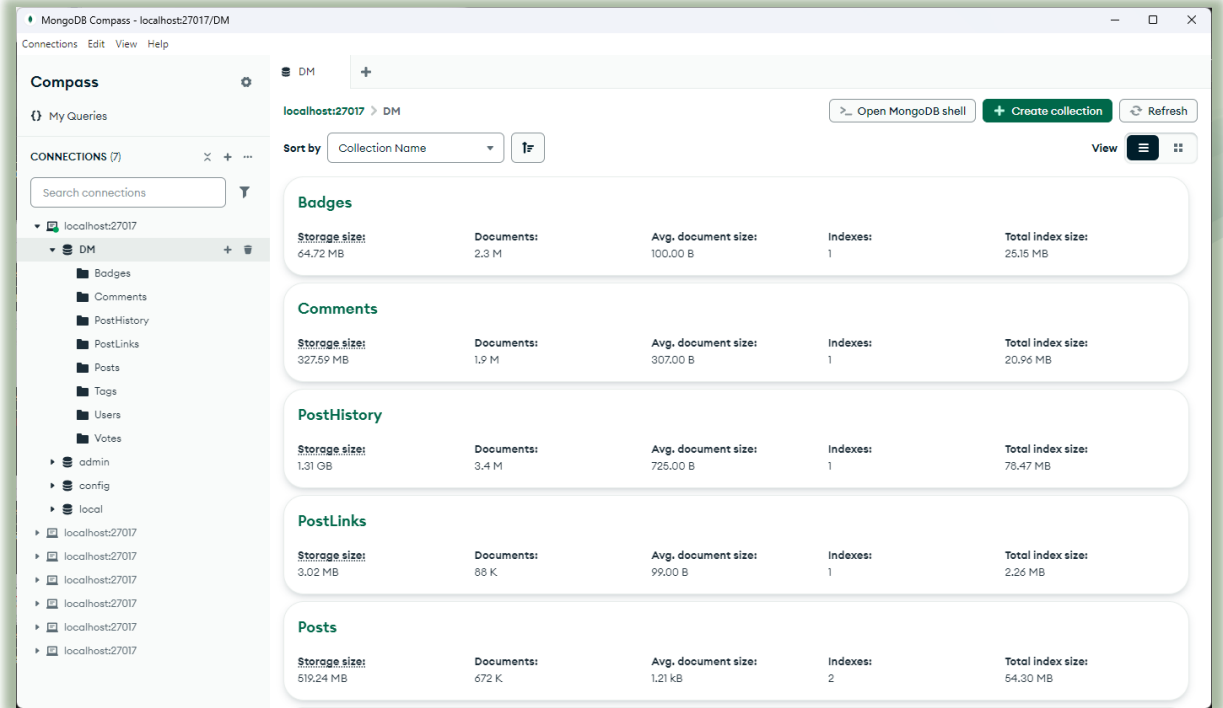


Interface Usability Layout

pgAdmin



MongoDB Compass

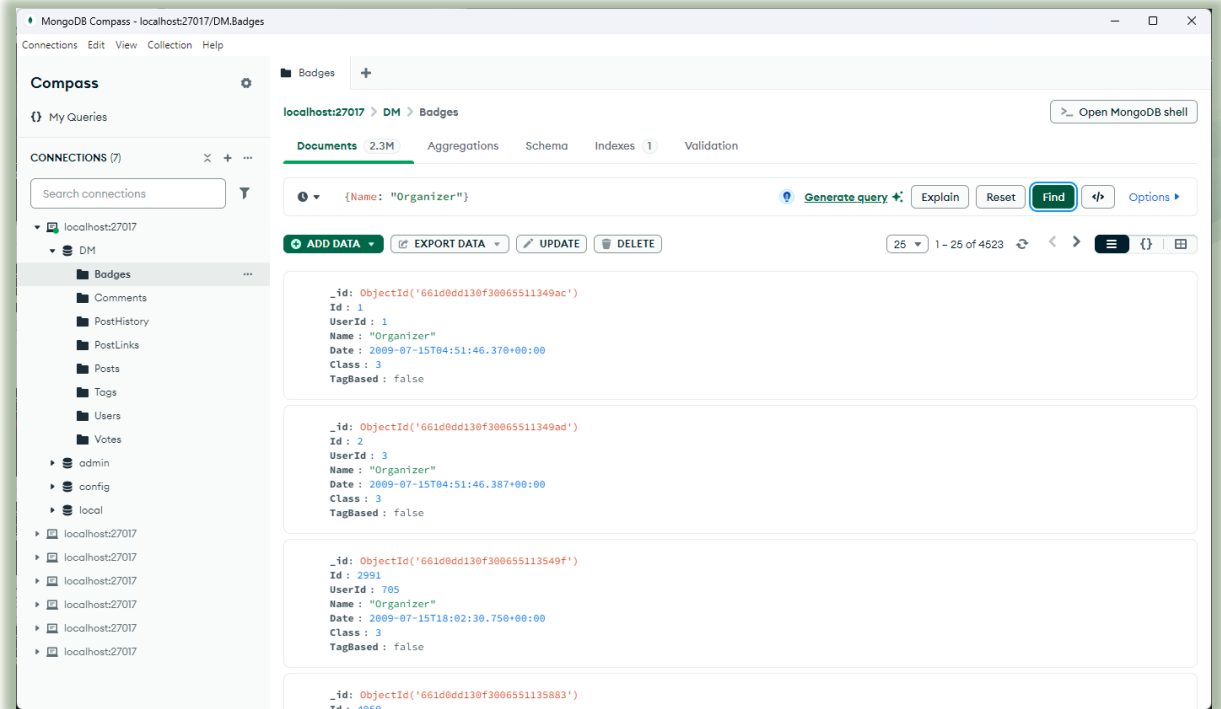
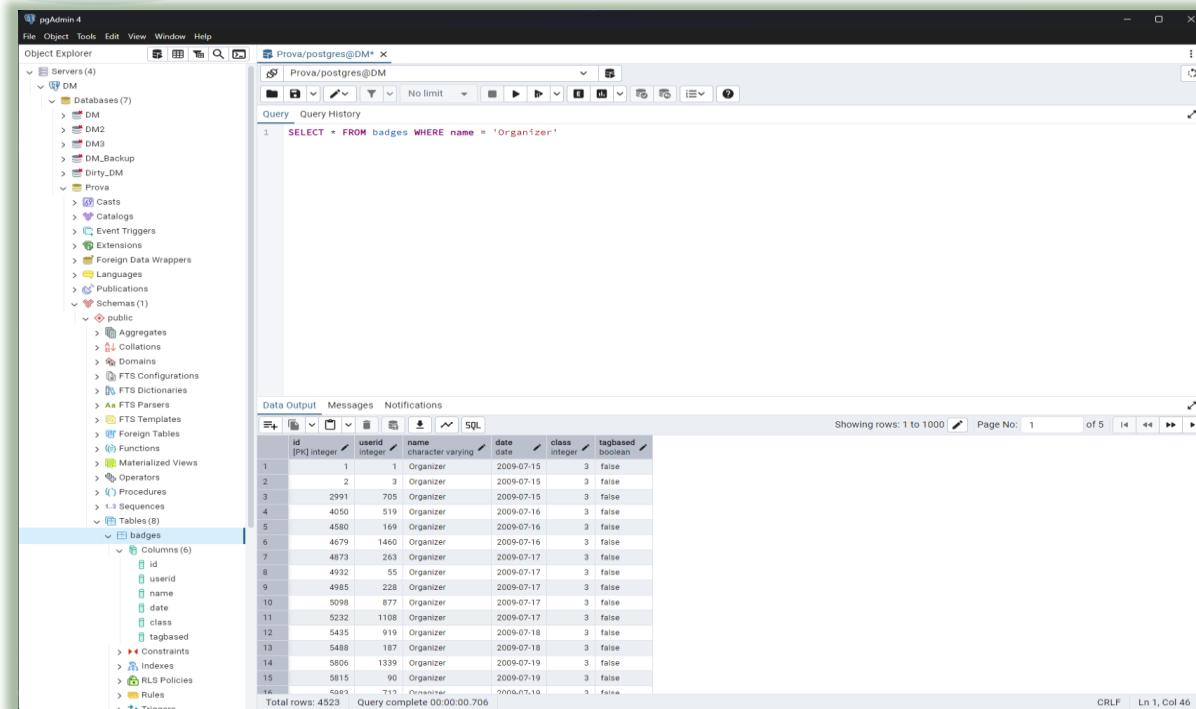


Interface Usability

Navigation & Usability

pgAdmin

MongoDB Compass

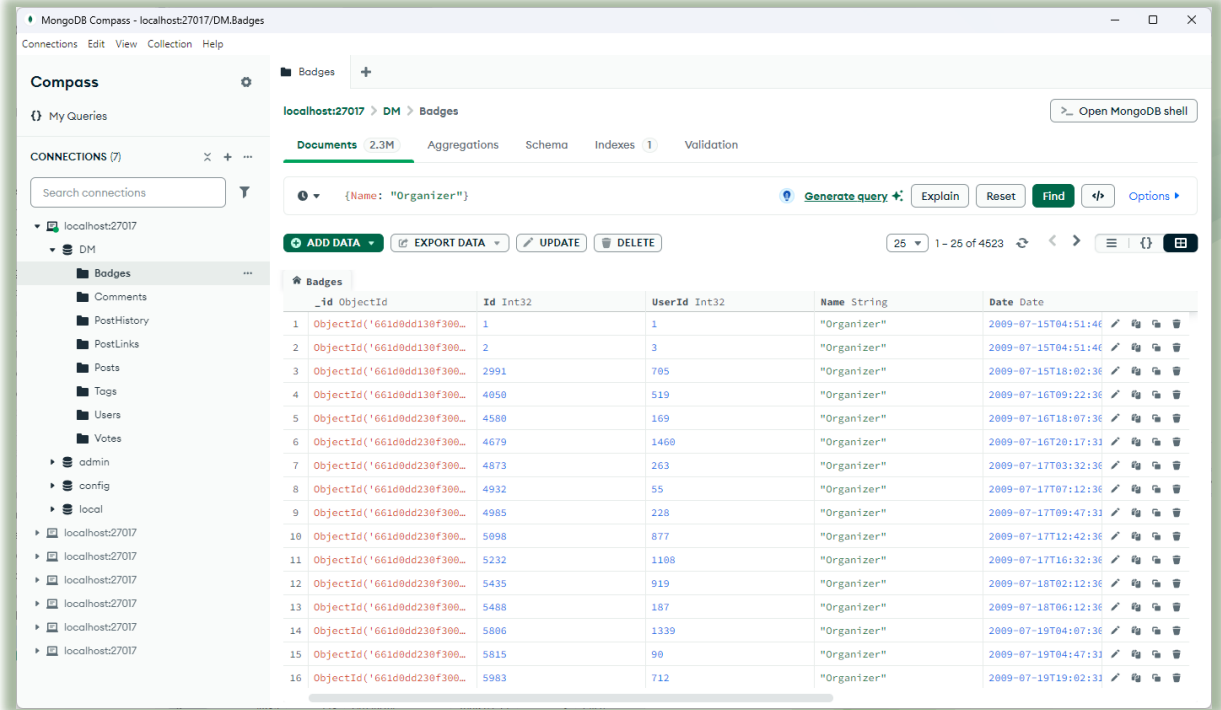
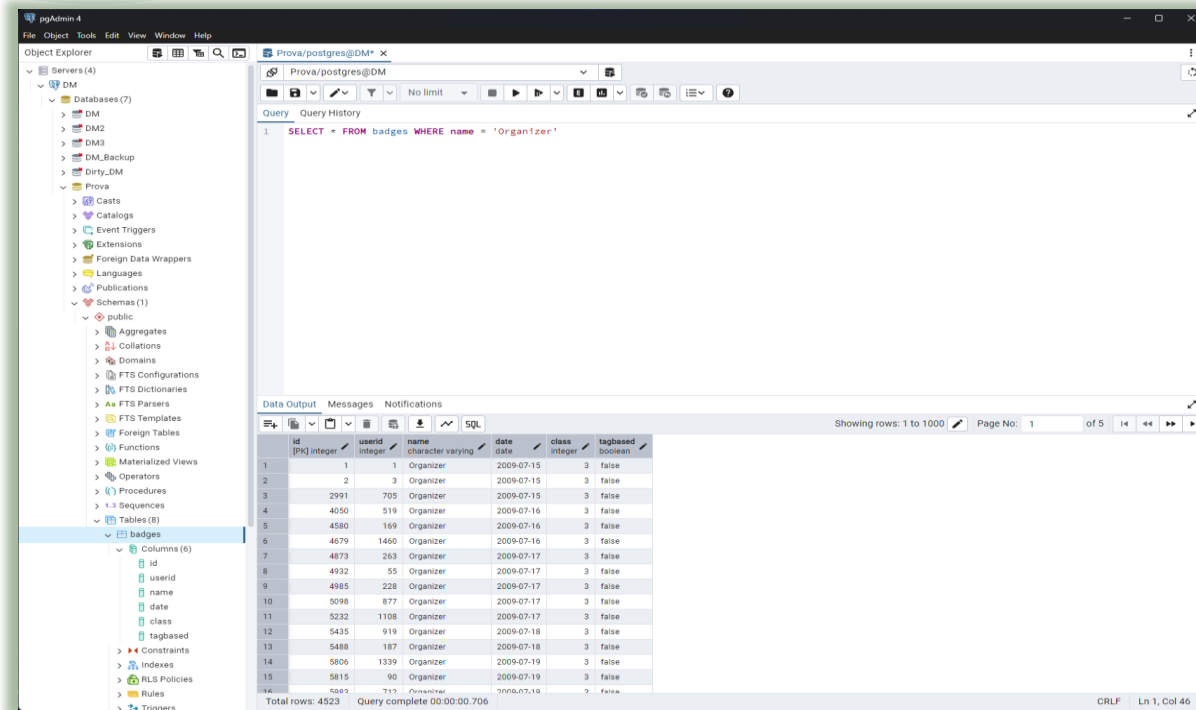


Interface Usability

Navigation & Usability

pgAdmin

MongoDB Compass



Interface Usability

Functionality

➤ Differences

Query Interface:

- pgAdmin: Uses a simple text editor for writing SQL queries.
- MongoDB Compass: Provides a visual aggregation pipeline builder for NoSQL queries.

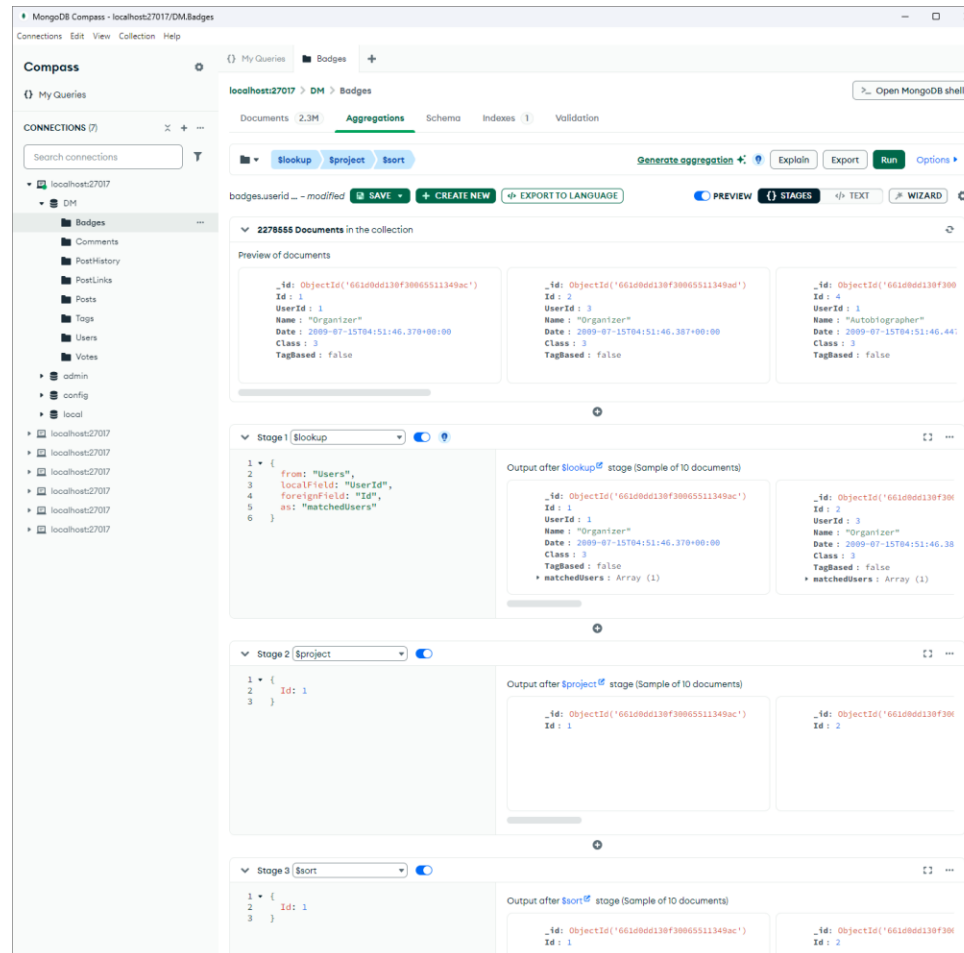
Predefined Queries:

- pgAdmin: Offers predefined queries (e.g., view data, CRUD scripts) via right-click.
- MongoDB Compass: No predefined SQL-like queries, relies on document-based query building.

➤ Similarities

- Both provide a GUI interface for database management.
- Both allow import/export of data.
- Both include a CLI window for command-line interaction.

Interface Usability Functionality



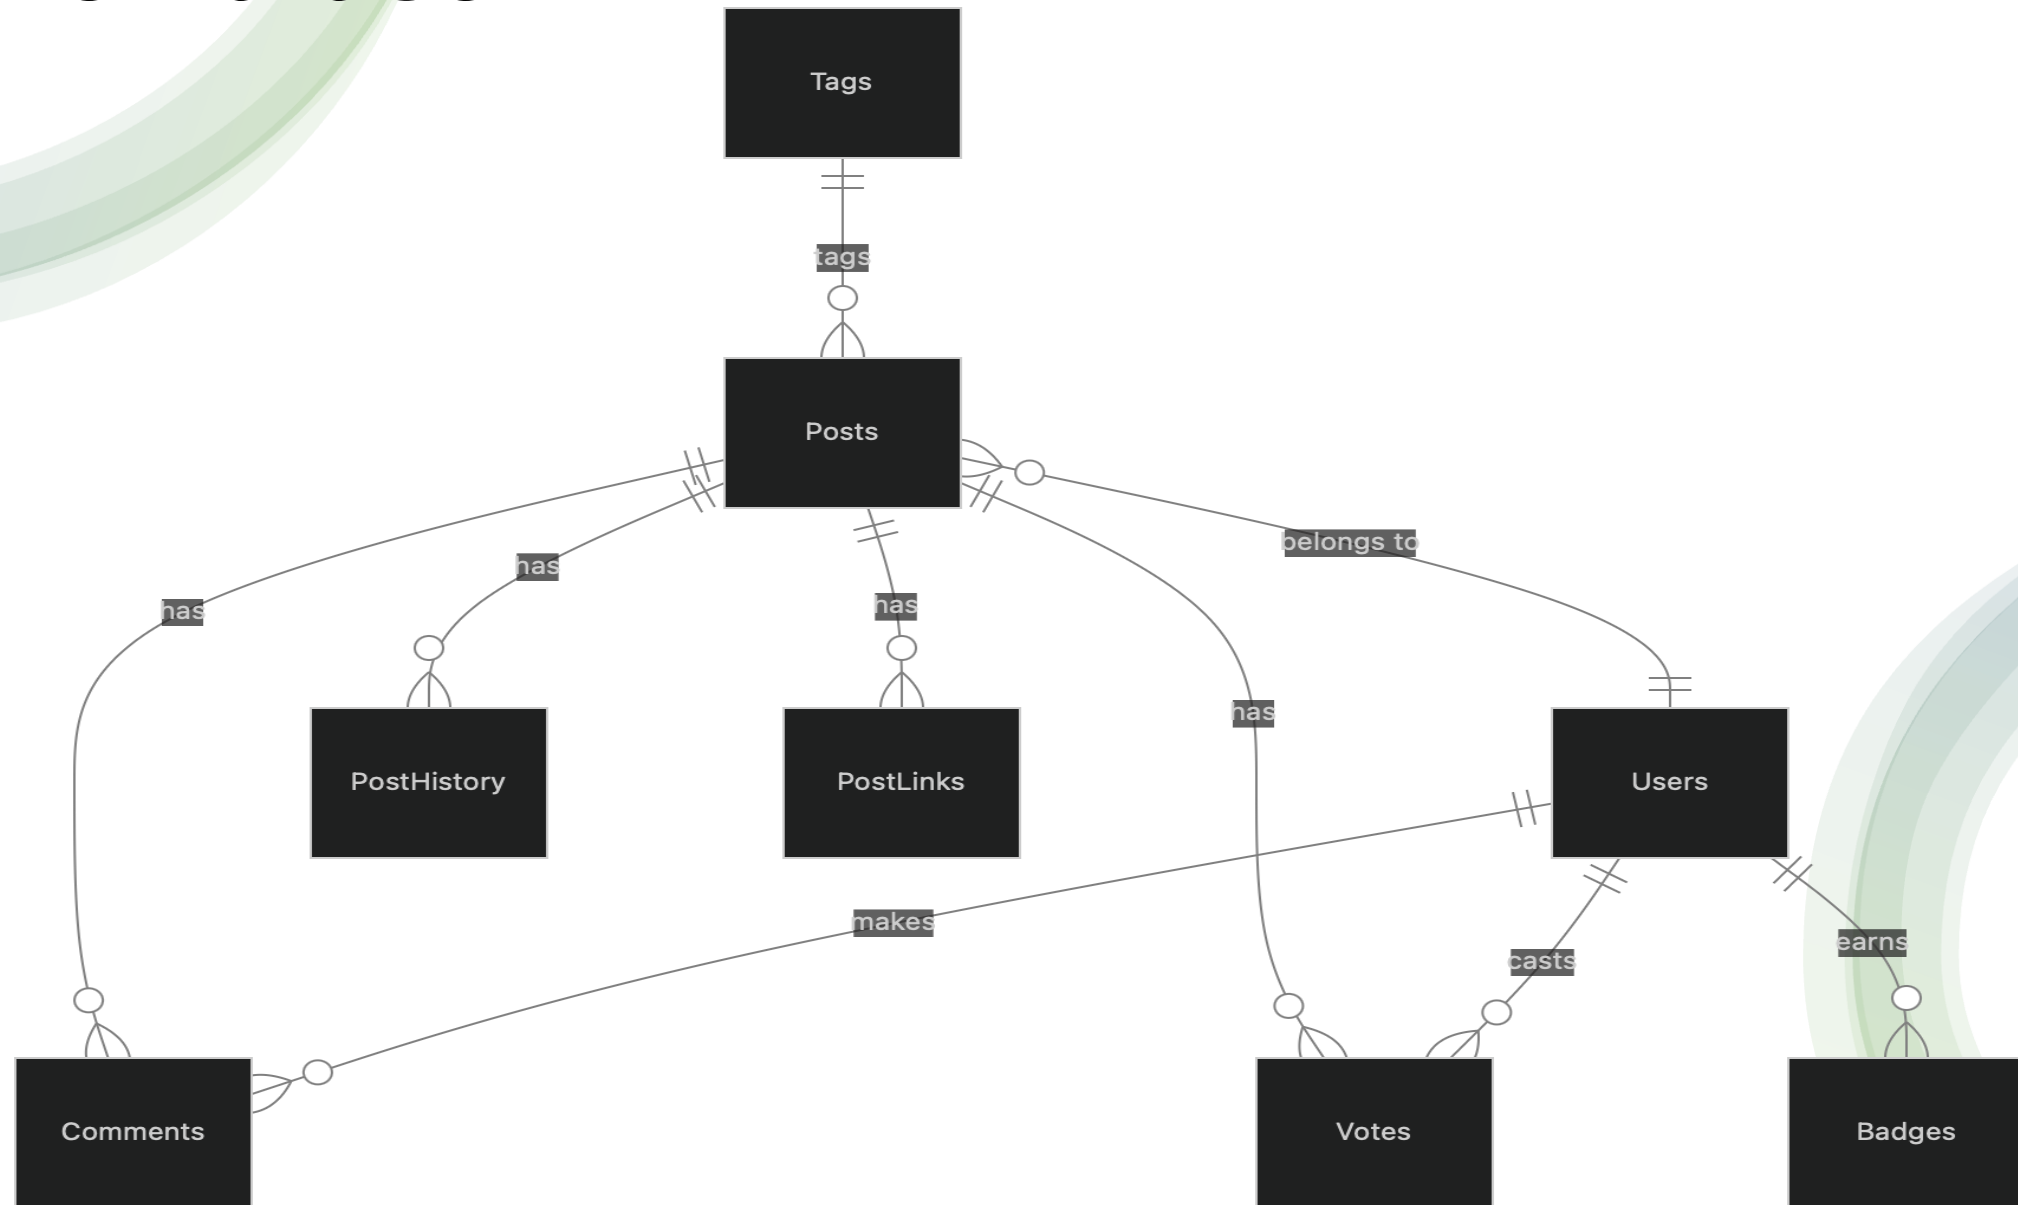
Flexibility

Feature	SQL (Relational Databases)	MongoDB (NoSQL)
Schema	Strict, predefined schema	Schema-less (Flexible documents)
Structure	Tables, rows, and columns	Documents (JSON/BSON format)
Data Relationships	Normalized, uses joins to relate tables	Denormalized, embeds related data
Schema Changes	Requires migrations, downtime	Dynamic, fields can be added/removed
Use Case	Best for structured & transactional data	Best for unstructured, evolving data

Scalability

Feature	SQL (Relational Databases)	MongoDB (NoSQL)
Scaling Strategy	Vertical Scaling (Scale-Up)	Horizontal Scaling (Scale-Out)
Data Distribution	Difficult to shard, manual partitioning	Built-in sharding for automatic distribution
Read/Write Performance	Efficient for transactions, joins slow large queries	Faster reads, distributed writes
Replication	Master-slave or cluster-based replication	Replica sets with automatic failover
High Availability	Requires complex setup	Native support for replication & redundancy
Global Distribution	Requires additional tools	Geo-distributed databases supported

The Dataset



Importing the tables

➤ PostgreSQL:

- Badges: 12,78 sec
- Comments: 16,87 sec
- PostHistory: 62 sec
- PostLinks: 0,72 sec
- Posts: 27,38 sec
- Tags: 0,11 sec
- Users: 14,72 sec
- Votes: 23,34 sec

➤ MongoDB doesn't measure time

Performance



Integrity Checks



Data Cleaning



Foreign Keys
(only for SQL)



Different Query
Types

Key Observation

Checking if PostLinks references Posts

- **Inefficient query:**

```
SELECT pl.id
FROM postlinks pl
WHERE pl.postid NOT IN (SELECT p.id FROM posts p);
```

Execution Time: 2 hours 37 minutes

- **Efficient query:**

```
SELECT pl.id
FROM postlinks pl
EXCEPT
SELECT pl.id
FROM postlinks pl JOIN posts p ON p.id = pl.postid
ORDER BY id;
```

Execution Time: 258 milliseconds

Key Observation #2

- MongoDB requires some indexing for very large data
- The queries couldn't complete/took very long without them
- The following indexes have been created:
 - Posts(Id)
 - Users(Id)

Integrity Checks

Among 8 queries	PostgreSQL	MongoDB
BEST PostLinks → Posts	258 ms	1.825 sec
AVERAGE	4.212 sec	33.466 sec (~7x)
WORST Votes → Posts	7.914 sec	76.631 sec

Data Cleaning

Among 8 queries	PostgreSQL	MongoDB
BEST PostLinks → Posts	379 ms	1.460 sec
AVERAGE	9.644 sec	34.108 sec (~2.5x)
WORST	29.958 sec PostHistory → Posts	87.51 sec Votes → Posts

- Votes table is **too large** for MongoDB's \$lookup.
- PostgreSQL's join is optimized on large tables.

Foreign Keys

➤ **Posts** references:

- Users(Id) → OwnerUserId
- Posts(Id) → ParentId & AcceptedAnswerId

➤ **Comments** references:

- Posts(Id) → PostId

➤ **Votes** references:

- Users(Id) → UserId
- Posts(Id) → PostId

➤ **Badges** references:

- Users(Id) → UserId

➤ **PostHistory** references:

- Users(Id) → UserId
- Posts(Id) → PostId

➤ **PostLinks** references:

- Posts(Id) → PostId & RelatedPostId

Performance Queries Overview

	Read	Write	Joins	Aggregation	Text Search	Nested Updates	Complex Joins	Pagination	Deletion	Count
SQL	0.779	0.453	14.673	2.236	1.286	3.259	8.998	0.583	>8h	0.884
MongoDB	6.867	0.003	88.112	2.518	2.3296	1.7278	>6h	2.470	10.869	1.144

Simple Read and Write

```
SELECT *  
FROM Posts  
WHERE CreationDate BETWEEN '2009-01-01' AND '2009-12-31';
```

- Declarative syntax
- Standard SQL

	Read	Write
SQL	0.779	0.453

```
INSERT INTO Posts (Id, PostTypeId, CreationDate, Score, Title, Body, Tags)  
VALUES  
(2, 1, NOW(), 10, 'Title 1', 'Body 1', '<mac><crash>'),  
(3, 2, NOW(), 15, 'Title 2', 'Body 2', '<windows>');
```

```
posts.find({"CreationDate":  
           {"$gte": datetime(2009, 1, 1), "$lte":  
datetime(2009, 12, 31)}})
```

- JSON-like documents
- Operators and expressions: \$gt, \$eq, ...

	Read	Write
MongoDB	6.867	0.003

```
posts.insert_many([  
{"Id": 2, "PostTypeId": 1, "CreationDate": datetime.now(), "Score": 10,  
"Title": "Title 1", "Body": "Body 1", "Tags": "<mac><crash>"},  
{"Id": 3, "PostTypeId": 2, "CreationDate": datetime.now(), "Score": 15,  
"Title": "Title 2", "Body": "Body 2", "Tags": "<windows>"  
})
```

Simple Joins and Aggregations

```
SELECT p.*, u.DisplayName  
FROM Posts p  
JOIN Users u ON p.OwnerUserId = u.Id;
```

- Optimized for joins
- Relationships are optimized

	Joins	Aggregation
SQL	14.673	2.236

```
SELECT Tags, COUNT(*)  
FROM Posts  
GROUP BY Tags;
```

- Aggregation with group by

```
posts.aggregate([  
  {"$lookup": {"from": "Users", "localField": "OwnerUserId",  
    "foreignField": "Id", "as": "owner"}},  
  {"$unwind": "$owner"},  
  {"$project": {"_id": 0, "Title": 1, "OwnerDisplayName":  
    "$owner.DisplayName"}}  
])
```

- No native join operation
- \$lookup stage is used in aggregations

	Joins	Aggregation
MongoDB	88.112	2.518

```
posts.aggregate([  
  {"$unwind": "$Tags"},  
  {"$group": {"_id": "$Tags", "count": {"$sum": 1}}}  
])
```

- Aggregation in stages

Text Search and Nested Updates

```
SELECT *  
FROM Posts  
WHERE Body LIKE '%virtual machine%';
```

	Text Search	Nested Updates
SQL	1.286	3.259

```
UPDATE Users  
SET Reputation = Reputation + 10  
WHERE Id IN  
(SELECT OwnerUserId  
FROM Posts  
GROUP BY OwnerUserId  
HAVING COUNT(*) > 10);
```

```
posts.find({"Body":  
{"$regex": "virtual machine", "$options": "i"}  
})
```

	Text Search	Nested Updates
MongoDB	2.3296	1.7278

```
owners_with_more_than_10_posts =  
posts.aggregate([  
{"$group": {"_id": "$OwnerUserId", "postCount": {"$sum": 1}}},  
{"$match": {"postCount": {"$gt": 10}}  
])  
owner_ids = [owner["_id"] for owner in  
owners_with_more_than_10_posts]  
users.update_many({"Id": {"$in": owner_ids}}, {"$inc": {"Reputation":  
10}})
```

Complex Joins

```
SELECT p.Title, c.Text, u.DisplayName
FROM Posts p
JOIN Comments c ON p.Id = c.PostId
JOIN Users u ON c.UserId = u.Id;
```

```
posts.aggregate([
  {"$lookup": {"from": "Comments", "localField": "Id", "foreignField": "PostId", "as": "comments"}},
  {"$unwind": "$comments"},
  {"$lookup": {"from": "Users", "localField": "comments.UserId", "foreignField": "Id", "as": "commentUser"}},
  {"$unwind": "$commentUser"},
  {"$project": {"_id": 0, "Title": 1, "CommentText": "$comments.Text", "CommenterName": "$commentUser.DisplayName"}}
])
```

	Complex Joins
SQL	8.998
MongoDB	>6h

Pagination, Deletion and Count

	Pagination	Deletion	Count
SQL	0.583	>8h	0.884
MongoDB	2.470	10.869	1.144

```
SELECT *  
FROM Posts  
ORDER BY CreationDate DESC LIMIT 10  
OFFSET 20;
```

```
DELETE FROM Posts  
WHERE Score < 5 AND Id NOT IN (SELECT  
PostId FROM Comments);
```

```
SELECT COUNT(*)  
FROM Users  
WHERE Reputation > 100;
```

```
posts.find().sort("CreationDate", -  
1).skip(20).limit(10)
```

```
post_ids_with_comments = comments.distinct("PostId")  
posts.delete_many({"Score": {"$lt": 5}, "Id": {"$nin":  
post_ids_with_comments}})
```

```
users.count_documents({"Reputation": {"$gt": 100}})
```

Conclusions

	Read	Write	Joins	Aggregation	Text Search	Nested Updates	Complex Joins	Pagination	Deletion	Count
SQL	0.779	0.453	14.673	2.236	1.286	3.259	8.998	0.583	>8h	0.884
MongoDB	6.867	0.003	88.112	2.518	2.3296	1.7278	>6h	2.470	10.869	1.144

- MongoDB better suited for changing data (faster write, update and delete operations)
- PostgreSQL faster in structured data

Use case of this dataset: large forum

- Dynamic data, but frequent join requests (loading posts and comments)
- Optimizations are put in place (e.g., on delete set null)



Thank you

Sahar Khanlari - 2107563

Marco Natale – 1929854

A.A. 2024/2025

